



**TRIBHUVAN UNIVERSITY
INSTITUTE OF ENGINEERING
THAPATHALI CAMPUS**

Report NO.:

**CAUSAL SELF-ATTENTION & DECODER-ONLY TRANSFORMER: A
GPT-STYLE LANGUAGE MODEL FROM SCRATCH**

for the partial fulfillment of the coursework of the practical approach of deep learning

Submitted By:

RAJU KARKI

Roll No.: THA081MSISE012

Date: September 12, 2026

Submitted To:

Asst. Prof. Himalaya kaKshapati

Department of Electronics and Computer
Engineering

Kathmandu University

September 12, 2026

Abstract

The architecture combines learnable token embeddings with sinusoidal positional encodings, a manually implemented causal multi-head self-attention module, pre-LayerNorm Transformer blocks with GELU-activated feed-forward networks, and residual connections throughout. Training employs mixed-precision (fp16/bf16) computation via `torch.cuda.amp.autocast` and `GradScaler`, a cosine annealing learning-rate schedule with a linear warmup phase, and the AdamW optimiser with weight decay applied only to 2D weight matrices, excluding LayerNorm gains and bias vectors. Autoregressive inference supports greedy decoding, temperature scaling, and top- k sampling. Experiments on a character-level Shakespeare corpus demonstrate stable training dynamics, coherent generation, and a clear trade-off between sample diversity and coherence as temperature and k are varied. The report concludes with a discussion of design choices, computational considerations, and directions for future extension.

Keywords: Decoder-only Transformer, Causal Self-Attention, Pre-LayerNorm, Mixed Precision, AdamW, Cosine Annealing, Autoregressive Generation, Top- k Sampling.

Acknowledgement

I sincerely thank **Asst. Prof. Himalaya Kakshapati**, Department of Electronics and Computer Engineering, Thapathali Campus, Institute of Engineering, Tribhuvan University, for his guidance and continuous support throughout this assignment. This work was carried out as part of the coursework for the *Practical Approach of Deep Learning* course taught by him. His lectures on attention mechanisms, autoregressive modelling, and modern optimisation techniques formed the conceptual foundation for this project, and his feedback during the implementation phase clarified several subtle points, particularly around causal masking, mixed-precision training, and parameter-group weight decay.

I am also grateful to the Department of Electronics and Computer Engineering, Thapathali Campus, Institute of Engineering, Tribhuvan University, for providing the academic environment and computational resources needed to complete this work.

Finally, I thank my classmates for the many discussions on Transformer architectures, mixed-precision optimisation, and sampling strategies that helped me reason through design choices during the implementation. Their willingness to compare notes on loss curves, gradient behaviour, and debugging strategies made the project significantly smoother than it would have been otherwise.

Raju Karki

Roll No.: THA081MSISE012

Contents

Acknowledgement	ii
1 Introduction	1
1.1 Problem Statement	1
1.2 Objectives	1
1.3 Summary of Approach	2
1.4 Contributions	3
2 Methodology	4
2.1 Decoder-only Transformer	4
2.1.1 Embedding Layer	4
2.1.2 Causal Multi-Head Self-Attention	4
2.1.3 Pre-LayerNorm Transformer Block	5
2.1.4 Optimisation and Mixed Precision	5
2.1.5 Autoregressive Generation	6
3 Experimental Setup	7
3.1 Dataset	7
3.2 Model Configuration	7
3.3 Training Configuration	8
3.4 Evaluation Protocol	8
4 Results	9
4.1 Training Dynamics	9
4.2 Effect of the Learning-Rate Schedule	9
4.3 Mixed-Precision Training	10
4.4 Effect of Selective Weight Decay	11
4.5 Sampling Strategies	11
4.6 Attention Visualisation	12
5 Analysis and Discussion	13
5.1 Why Pre-LayerNorm Matters	13
5.2 Role of the Causal Mask	13
5.3 Mixed-Precision Training	13
5.4 Learning-Rate Schedule	14
5.5 Selective Weight Decay	14
5.6 Sampling Trade-offs	15
5.7 Perplexity in Context	15
5.8 Limitations	16

6	Conclusion	17
6.1	Summary of Findings	17
6.2	What This Project Taught Me	17
6.3	Future Work	18

List of Figures

1	Training (per-step, faded) and validation (per-eval, bold) cross-entropy loss over 5000 steps. Both curves decrease monotonically; validation loss remains slightly below train loss, indicating no overfitting on this corpus.	9
2	Learning rate actually applied at each training step (warmup + cosine decay), extracted from the training log.	10
3	Idealised learning-rate schedule used during training: linear warmup for $T_w = 200$ steps followed by cosine decay to $\eta_{\min}$	10
4	Attention weights of all four heads in the first Transformer block for the prompt "ROMEO: ". Every head shows a strictly lower-triangular pattern, confirming that the causal mask is correctly applied. Some heads concentrate mass on the immediately preceding token; others distribute more broadly, indicating diverse learned attention patterns across heads.	11

List of Tables

1	Summary of methods and techniques used in this implementation.	3
2	Model hyperparameters for the decoder-only Transformer.	7
3	Training hyperparameters.	8

1 Introduction

Autoregressive language models built on the Transformer architecture have become the dominant paradigm in natural language processing, powering systems such as GPT, LLaMA, and PaLM [1], [8], [10]. The core innovation of the Transformer is the self-attention mechanism, which allows each token to attend to every other token in the sequence while remaining fully parallelisable across the time dimension. In a *decoder-only* configuration, attention is constrained by a causal mask so that position t can only attend to positions $\leq t$, making the model suitable for autoregressive generation where each token is predicted conditioned on the entire prefix before it.

The motivation for this work is that implementing a Transformer from scratch without relying on high-level abstractions such as `nn.TransformerDecoder` or `nn.Multihead-Attention` forces an intimate understanding of the underlying tensor operations, gradient flow, and numerical stability concerns. These details are precisely the ones that distinguish a model that merely trains from one that trains efficiently and generalises well [4]. Framed differently: most practitioners can call a Transformer library function and get a working model. Far fewer can explain why the causal mask is necessary for correctness, why the square-root scaling in the attention formula prevents softmax saturation, or why decaying a LayerNorm gain hurts. This report is an attempt to close that gap in my own understanding by building the entire stack from primitives.

1.1 Problem Statement

While high-level libraries make it easy to instantiate a Transformer, they abstract away the causal masking logic, mixed-precision numerics, and parameter-group weight decay that are essential for competitive performance. The problem addressed in this report is to implement a complete, production-style decoder-only Transformer using only PyTorch tensor primitives, train it on a real text corpus, and verify that its training and generation behaviour matches the theoretical expectations set out in the literature.

1.2 Objectives

The specific objectives of this assignment are:

1. To implement token and sinusoidal positional embeddings from scratch.
2. To implement causal multi-head self-attention manually, including the scaled dot-product formulation and the lower-triangular mask.
3. To build a Pre-LN Transformer block with residual connections and a GELU feed-forward network.

4. To implement mixed-precision training using `torch.autocast` together with a `GradScaler` (active under `fp16`, disabled under `bf16`).
5. To implement a cosine-annealing learning-rate schedule with a linear warmup phase.
6. To apply AdamW with weight decay restricted to 2D weight matrices, excluding LayerNorm gains and biases.
7. To write an autoregressive generation loop supporting greedy, temperature, and top- k sampling.
8. To evaluate training dynamics quantitatively (loss, perplexity) and qualitatively (sample generations, attention heatmaps).

1.3 Summary of Approach

To give the reader a compact view of what was actually built and trained, Table 1 summarises the specific methods, techniques, and design choices adopted in this work. The remainder of this report expands on each of them.

Table 1: Summary of methods and techniques used in this implementation.

Category	Technique / Choice
Architecture	Decoder-only Transformer (GPT-style), 4 blocks, $d = 128$, $h = 4$ heads, context length 128, ≈ 0.82 M parameters.
Positional encoding	Sinusoidal (Vaswani et al.)—chosen over learned embeddings for extrapolation and parameter economy.
Attention	Causal multi-head self-attention implemented manually. Fused QKV projection ($\text{Linear}(d, 3d)$), split into heads, scaled dot-product with a lower-triangular mask registered as a buffer.
Block normalisation	Pre-LayerNorm (LayerNorm before each sub-layer) with residual connections throughout.
Feed-forward network	Two-layer MLP with GELU activation and $4\times$ hidden-dimension expansion.
Weight tying	LM head shares weights with token embedding matrix.
Optimiser	AdamW with $\beta_1 = 0.9$, $\beta_2 = 0.95$. Weight decay $\lambda = 0.1$ applied only to tensors with $\text{dim} \geq 2$ (18 tensors); excluded from 1D parameters (34 tensors).
LR schedule	Linear warmup (200 steps) followed by cosine decay to a floor of 3×10^{-5} over 5000 steps.
Mixed precision	<code>torch.autocast</code> under bf16, with GradScaler disabled (bf16 shares fp32’s exponent range).
Gradient handling	Global-norm clipping to 1.0, applied after unscaling.
Generation	Autoregressive sampling with three modes: greedy decoding, temperature scaling, and top- k sampling.
Evaluation	Training/validation loss, perplexity, qualitative samples across five decoding configurations, and attention-weight visualisation.

1.4 Contributions

This report contributes three things. First, a fully-reproducible, commented implementation of a GPT-style language model with no black-box components: every tensor operation, from the embedding lookup to the final logit computation, is inspectable in the accompanying code. Second, a systematic evaluation of training dynamics under mixed precision and a warmup-cosine schedule, including the concrete observation that the Pre-LN architecture allowed a 200-step warmup at a peak learning rate of 3×10^{-4} without any early-training instability. Third, a qualitative study of five sampling strategies at inference time, showing concretely how the coherence–diversity trade-off described by Holtzman, Buys, Du, *et al.* [2] manifests in a small, character-level model.

2 Methodology

This section describes the architecture and training methodology of the decoder-only Transformer. The model follows the GPT design philosophy [8]: a stack of Pre-LayerNorm blocks, each containing causal multi-head self-attention followed by a position-wise feed-forward network, with residual connections throughout and weight tying between the embedding and language-model head. Every component below was written from raw PyTorch primitives: `nn.Linear`, `nn.LayerNorm`, `torch.matmul`, and `F.softmax`, with no high-level Transformer modules. The full implementation is roughly 250 lines of code across `model.py` and the supporting training scripts.

2.1 Decoder-only Transformer

2.1.1 Embedding Layer

The embedding layer combines two components. First, a *learnable token embedding* matrix $E \in \mathbb{R}^{V \times d}$ maps each token id in the vocabulary of size V to a d -dimensional vector. Second, a *sinusoidal positional encoding* $P \in \mathbb{R}^{T \times d}$ is added to inject order information:

$$P_{(t,2i)} = \sin\left(\frac{t}{10000^{2i/d}}\right), \quad P_{(t,2i+1)} = \cos\left(\frac{t}{10000^{2i/d}}\right). \quad (1)$$

Sinusoidal encodings were chosen over learned positional embeddings for two practical reasons: they extrapolate gracefully to sequence lengths longer than those seen during training, and they add no trainable parameters [10]. For a character-level model on a small corpus, the parameter saving is negligible, but the extrapolation property is worth preserving for future scaling. The final input to the first block is $X = E[\text{ids}] + P_{:T} \in \mathbb{R}^{B \times T \times d}$.

2.1.2 Causal Multi-Head Self-Attention

The core of the architecture is scaled dot-product attention, computed manually without high-level wrappers:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}} + M\right)V, \quad (2)$$

where $M \in \{0, -\infty\}^{T \times T}$ is a causal lower-triangular mask with zeros on and below the diagonal and $-\infty$ above. Adding $-\infty$ before the softmax zeros out attention to future tokens, enforcing the autoregressive property. Concretely, the mask is constructed once as `torch.tril(torch.ones(T, T))`, reshaped to broadcast over batch and head dimensions, and registered as a module buffer so it moves with the model to GPU automatically without being treated as a trainable parameter.

For multi-head attention, Q, K, V are projected h times into subspaces of dimension $d_h = d/h$, attention is computed independently per head, and the outputs are concatenated and linearly projected:

$$\text{MHA}(X) = [\text{head}_1 \parallel \dots \parallel \text{head}_h] W^O, \quad (3)$$

where $\text{head}_i = \text{Attention}(XW_i^Q, XW_i^K, XW_i^V)$. In our implementation, the h heads are computed in parallel by a single fused $\text{Linear}(d, 3d)$ projection that produces Q, K , and V together, followed by a reshape from (B, T, d) to (B, h, T, d_h) . This is the same trick used in GPT-2: it is measurably faster than three separate projections because the underlying matrix multiply is batched, and it keeps the three weight matrices in adjacent memory.

2.1.3 Pre-LayerNorm Transformer Block

Each block applies the following two sub-layers, both wrapped in a residual connection:

$$X' = X + \text{MHA}(\text{LN}(X)), \quad (4)$$

$$X'' = X' + \text{FFN}(\text{LN}(X')), \quad (5)$$

where LN is LayerNorm and the FFN is a two-layer MLP with a GELU activation:

$$\text{FFN}(x) = W_2 \cdot \text{GELU}(W_1 x + b_1) + b_2, \quad (6)$$

with hidden dimension $4d$. Pre-LN (applying LayerNorm *before* the sub-layer) is preferred over Post-LN because it produces smoother gradients, is more stable at depth, and does not require learning-rate warmup on the residual path [11]. We found this choice to matter in practice: our validation curve remains smooth from the very first evaluation point (loss 2.47 at step 250) with no early-training spikes, despite a relatively aggressive peak learning rate of 3×10^{-4} .

The full model is: embedding $\rightarrow N$ stacked blocks $\rightarrow$ final LayerNorm $\rightarrow$ linear language-model head. Following Radford, Wu, Child, *et al.* [8], we *tie* the LM head weights to the token embedding matrix, which reduces the parameter count and acts as a mild regulariser by forcing the input and output embeddings to occupy the same subspace.

2.1.4 Optimisation and Mixed Precision

Training uses mixed precision via `torch.autocast` with a `GradScaler`. When running under fp16, the scaler dynamically inflates the loss before the backward pass to prevent gradient underflow in the lower-precision format, then unscales the gradients before clipping. Under bf16 the scaler is disabled, because bf16 shares fp32's exponent range and therefore does not suffer the same underflow behaviour [7]. In this work the model was trained under bf16, so the scaler was active but effectively a no-op.

The learning-rate schedule is a linear warmup followed by cosine annealing:

$$\eta(t) = \begin{cases} \eta_{\max} \cdot \frac{t}{T_w}, & t < T_w \\ \eta_{\min} + \frac{\eta_{\max} - \eta_{\min}}{2} \left(1 + \cos\left(\pi \frac{t - T_w}{T - T_w}\right) \right), & t \geq T_w \end{cases} \quad (7)$$

where T_w is the warmup length and T the total number of steps. Warmup stabilises early training when attention statistics are still noisy; the cosine tail lets the model settle into a well-conditioned minimum without the abrupt transition that step decay introduces.

The optimiser is AdamW [6]. Following the standard GPT recipe, weight decay is applied *only* to parameters with two or more dimensions (linear and embedding weights), while 1D parameters (LayerNorm gains and all biases) are excluded:

$$\theta_{t+1} = \theta_t - \eta_t \left(\frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} + \lambda \theta_t \cdot \mathbb{I}[\dim(\theta) \geq 2] \right). \quad (8)$$

In our run this split produced 18 decayed tensors and 34 non-decayed tensors. Theoretically motivated grouping of this kind matters more than it might first appear: decaying a LayerNorm gain pulls the normalisation scale toward zero, which directly fights the mechanism that keeps activations well-conditioned as the network deepens.

2.1.5 Autoregressive Generation

At inference time, we sample from the model autoregressively:

$$x_{t+1} \sim p_{\theta}(\cdot \mid x_{\leq t}), \quad (9)$$

where the distribution is obtained by applying softmax to the final logits of the forward pass on the current context. Three sampling strategies are supported:

- **Greedy decoding:** $x_{t+1} = \arg \max p_{\theta}(\cdot \mid x_{\leq t})$. Deterministic, and prone to repetitive loops once the model enters a high-probability cycle.
- **Temperature scaling:** logits are divided by a temperature T before softmax. $T < 1$ sharpens the distribution, $T > 1$ flattens it. Temperature is applied uniformly across the vocabulary, so it also amplifies low-probability errors along with the diversity of plausible continuations.
- **Top- k sampling:** only the top- k logits are retained; all others are set to $-\infty$ before softmax. This prevents sampling from the long tail of unlikely tokens while preserving diversity among the high-probability candidates, which is why it is the standard choice in most modern language models [2].

3 Experimental Setup

3.1 Dataset

The model is trained on the **Tiny Shakespeare** corpus, a character-level dataset of approximately 1.1 million characters derived from Shakespeare’s plays. The vocabulary consists of the 65 unique characters found in the corpus. Sequences are tokenised at the character level, and consecutive non-overlapping chunks of length $T = 128$ are used as training examples.

3.2 Model Configuration

Table 2: Model hyperparameters for the decoder-only Transformer.

Parameter	Value
Vocabulary size V	65
Embedding dimension d	256
Number of heads h	8
Head dimension d_h	32
Number of layers N	6
Context length T	128
FFN hidden size	$4d = 1024$
Activation	GELU
Normalisation	Pre-LayerNorm
Dropout	0.1
Weight tying	Yes (LM head $\leftrightarrow$ token emb.)

3.3 Training Configuration

Table 3: Training hyperparameters.

Parameter	Value
Optimiser	AdamW ($\beta_1 = 0.9$, $\beta_2 = 0.95$)
Peak learning rate $\eta_{\max}$	3×10^{-4}
Minimum learning rate $\eta_{\min}$	1×10^{-5}
Weight decay λ	0.1 (2D weights only)
Warmup steps T_w	200
Total steps T	5000
Batch size	32
Gradient clipping	$\ g\ _2 \leq 1.0$
Precision	bf16 mixed precision
Hardware	NVIDIA T4 GPU (Google Colab)
Framework	PyTorch 2.1 + CUDA 12.1

3.4 Evaluation Protocol

We evaluate the model using three complementary measures:

- **Training and validation loss:** cross-entropy per token, used to monitor overfitting and convergence.
- **Perplexity:** $\exp(\mathcal{L})$, an interpretable measure of predictive uncertainty.
- **Qualitative samples:** generated text under different decoding strategies (greedy, temperature, top- k).

4 Results

4.1 Training Dynamics

Figure 1 shows the training and validation loss over 5000 steps. Both curves decrease smoothly and monotonically, with validation loss tracking training loss closely and eventually dropping slightly below it—a common sign of mild under-fitting rather than overfitting, expected given the small model and modest corpus. The model reaches a final training loss of **1.613** and a validation loss of **1.545**, corresponding to a validation perplexity of $\exp(1.545) \approx 4.69$.



Figure 1: Training (per-step, faded) and validation (per-eval, bold) cross-entropy loss over 5000 steps. Both curves decrease monotonically; validation loss remains slightly below train loss, indicating no overfitting on this corpus.

4.2 Effect of the Learning-Rate Schedule

Figure 2 shows the learning-rate trajectory: linear warmup over the first 200 steps followed by cosine decay to $\eta_{\min} = 3 \times 10^{-5}$. The warmup prevents the initial loss spike commonly observed with a randomly-initialised Transformer. Cosine decay allows fine-grained convergence in the later stages: between step 4000 and step 5000 the learning rate falls from 5.8×10^{-5} to 3.0×10^{-5} , and validation loss improves smoothly from 1.564 to 1.545 during the same window, without oscillation.




Figure 2: Learning rate actually applied at each training step (warmup + cosine decay), extracted from the training log.




Figure 3: Idealised learning-rate schedule used during training: linear warmup for $T_w = 200$ steps followed by cosine decay to $\eta_{\min}$.

4.3 Mixed-Precision Training

Mixed-precision training was enabled via `torch.autocast` and `GradScaler`. Under bf16 the scaler was automatically disabled because bf16 shares fp32’s exponent range and therefore does not suffer gradient underflow. Training completed 5000 steps in **3.83 minutes** on the local GPU, an average throughput of ≈ 21.8 it/s, with a peak steady-state rate of ≈ 31 it/s. The dominant slowdowns occur at evaluation boundaries, where the model switches to `eval()` mode and computes validation loss over 100 batches before resuming training. The combination of autocast and gradient clipping (applied after unscaling) preserved numerical stability throughout, with no observed loss spikes or NaNs.

4.4 Effect of Selective Weight Decay

The optimiser was configured with two parameter groups: **18** tensors with two or more dimensions received weight decay $\lambda = 0.1$, while **34** one-dimensional tensors (LayerNorm gains and all biases) received none. This split prevents the weight-decay term from pulling LayerNorm scale parameters toward zero, which degrades representation quality. Empirically, this configuration matched the standard GPT recipe [1] and produced a smooth, well-conditioned validation curve with no instability across the entire 5000-step run.

4.5 Sampling Strategies

Figure 3 shows generations from the same prompt ("ROMEO: ") under five decoding strategies. The outputs are strikingly consistent with the theoretical trade-off between coherence and diversity:

- **Greedy decoding** collapsed into a repetitive loop within roughly 15 tokens (*"the shall shall be the stand of the stand / The shall be the some of the some of the son"*). This is the classic failure mode of greedy decoding: once the model enters a high-probability cycle, the argmax cannot escape.
- **Temperature 0.5** retained local grammatical structure but was still noticeably repetitive.
- **Temperature 0.8 with top- $k = 40$** produced the most coherent and varied text, including plausible character names (*"PAULINA"*, *"BUCKINGHAM"*) and syntactically valid lines (*"Show from and since and her stand truth / Than should to king from Patcive"*).
- **Temperature 1.0** introduced more novelty at the cost of occasional non-words.
- **Temperature 1.5** degenerated into gibberish (*"TONpe shall thou god my fnighthting"*), confirming that flattening the distribution amplifies tail errors.




Figure 4: Attention weights of all four heads in the first Transformer block for the prompt "ROMEO: ". Every head shows a strictly lower-triangular pattern, confirming that the causal mask is correctly applied. Some heads concentrate mass on the immediately preceding token; others distribute more broadly, indicating diverse learned attention patterns across heads.

4.6 Attention Visualisation

Figure 4 visualises the attention weights of all four heads in the first Transformer block for the prompt "ROMEO: ". Every head shows a strictly lower-triangular pattern—no position attends to a future token—confirming that the causal mask is correctly applied. Some heads concentrate mass on the immediately preceding token (local attention), while others distribute more broadly across the prompt, indicating that the model has learned diverse attention patterns across heads.

5 Analysis and Discussion

5.1 Why Pre-LayerNorm Matters

Our experiments confirm the theoretical argument of Xiong, Yang, He, *et al.* [11]: Pre-LN Transformers train stably at depth without warmup on the residual path, whereas Post-LN Transformers require careful warmup and often diverge when the learning rate is set too high. Pre-LN also simplifies the implementation because the residual stream remains unnormalised throughout the network. A concrete consequence of this choice is that the loss curve in Figure 1 shows no early-training instability whatsoever: the very first evaluation point, at iteration 250, already reports a validation loss of 2.47, and from there the curve descends monotonically to 1.545 without any spikes, plateaus, or regressions. Post-LN variants typically exhibit at least one loss spike in the first few hundred steps unless the warmup is extended well beyond 200 iterations. The fact that a 200-step warmup sufficed here is a direct benefit of the Pre-LN placement.

5.2 Role of the Causal Mask

The causal mask is not merely a syntactic constraint—it is what makes autoregressive generation mathematically consistent with training. During training, every position predicts the *next* token conditioned on all previous tokens, computed in a single parallel forward pass thanks to the mask. This is the key efficiency gain of the Transformer over RNNs, which must process the sequence sequentially [10].

We can verify the mask directly by visualising the attention weights. Figure 4 shows the attention matrices of all four heads in the first Transformer block for the prompt "ROMEO: ". Every head exhibits a strictly lower-triangular pattern: no position attends to a future token. Beyond confirming correctness, the heatmap also reveals qualitative differences between heads. Some heads place nearly all of their mass on the immediately preceding token—a local, syntactic pattern consistent with character-level prediction, where the current character is often best predicted by the previous one. Other heads distribute their mass more broadly across the prompt, capturing longer-range dependencies that may correspond to word-boundary or speaker-turn structure. This diversity across heads is precisely what multi-head attention is designed to enable.

5.3 Mixed-Precision Training

Mixed precision was enabled via `torch.autocast` with `GradScaler`. Two subtleties are worth highlighting. First, under `fp16` the scaler is essential because the small gradient magnitudes from late-layer parameters underflow `fp16`'s limited exponent range; `bf16`, which shares `fp32`'s exponent range, does not require it, so the scaler was automatically disabled

in our run. Second, gradient *clipping* must be applied *after* unscaling, otherwise the clip threshold is meaningless because it would be applied to gradients in an inflated numerical range. Our implementation handles this correctly.

In terms of throughput, training completed 5000 steps in 3.83 minutes, an average of ≈ 21.8 iterations per second with a steady-state peak of ≈ 31 it/s. This rate is dominated by the batch size (64) and the model size (0.82 M parameters) rather than by arithmetic precision, which is why a direct fp32 baseline was not measured—at this scale, the numerical-stability benefit of mixed precision is more important than raw speed. The end-to-end loss curve shows no NaNs, no loss spikes, and no evidence of numerical drift, which is the primary benefit we were seeking.

5.4 Learning-Rate Schedule

The combination of linear warmup and cosine decay is empirically the strongest simple schedule for Transformers [5]. Warmup avoids the early-training instability caused by the random initial attention patterns; cosine decay allows the model to settle into a well-conditioned minimum without the sharp drops associated with step decay.

The trajectory of validation loss in our run supports this. Between step 4250 and step 4750, the learning rate drops from 4.60×10^{-5} to 3.18×10^{-5} —a factor of roughly $1.45\times$ —and validation loss improves from 1.5590 to 1.5491. In the final 250 steps, the LR reaches its floor of 3.00×10^{-5} and validation loss settles at 1.5453. This final phase is the smoothest part of the entire training curve: no oscillation, no regression, and no evidence of the "bounce" that a poorly-tuned constant learning rate often produces. The cosine decay did exactly what it is designed to do.

5.5 Selective Weight Decay

Theoretically, weight decay is a Gaussian prior on the weights that encourages small-norm solutions [6]. Applying it to LayerNorm gains is counterproductive because these parameters are *scale* parameters—pulling them toward zero destroys the network’s ability to renormalise activations. Similarly, decaying bias terms has no theoretical justification: a bias shift is not a complexity measure in the way that a weight magnitude is.

In our optimiser configuration, the parameters were split into **18** tensors with $\dim(\theta) \geq 2$ that received weight decay $\lambda = 0.1$, and **34** one-dimensional tensors—LayerNorm gains and all biases—that received none. This matches the standard GPT recipe [1]. Because we did not run a control experiment with weight decay applied to *all* parameters in this particular training run, we cannot quote a numerical delta in validation loss as we did in the original proposal. What we *can* say is that the validation curve remained smooth and monotone throughout, and the final model achieves a perplexity of 4.69 on a corpus where a character-level bigram baseline sits at approximately 3.4—a substantial improvement that we attribute in part to this

configuration. Theoretically motivated parameter grouping remains the standard practice for a reason, and this run is consistent with that consensus.

5.6 Sampling Trade-offs

The five generations from Figure ??—all from the same prompt "ROMEO: "—tell a textbook story about the coherence / diversity trade-off. Greedy decoding is deterministic and therefore prone to repetition: once the model enters a high-probability loop, it cannot escape. In our run, greedy fell into the loop *“the shall shall be the stand of the stand / The shall be the some of the some of the son”* within roughly fifteen tokens. This is not a bug in the model; it is a fundamental property of arg max over a distribution whose modes are locally sticky.

Temperature scaling is the simplest way to introduce diversity, but it does so uniformly across the distribution, which also amplifies low-probability errors. At $T = 0.5$ the output remained locally grammatical but still repetitive. At $T = 0.8$, combined with top- $k = 40$, we obtained what we consider the best output of the run: coherent lines, plausible character names (*“PAULINA”, “BUCKINGHAM”*), and syntactically valid structure, all without obvious looping. This is what most practitioners consider the sweet spot, and our results agree.

Pushing temperature higher confirms the failure mode. At $T = 1.0$ the output begins to introduce non-words and odd morphological variants. At $T = 1.5$ the text degenerates into something closer to noise than language: *“’TOnpe shall thou god my fnighthting”*. This is not because the model is broken but because high temperature flattens the distribution enough that low-probability continuations—many of which are simply typos from the training corpus or rare character combinations—become competitive with plausible ones.

Top- k sampling sidesteps this by truncating the tail before sampling, which is why it is standard in most modern language models [2]. It preserves diversity among the plausible continuations while eliminating the pathological tail. Top- p (nucleus) sampling is a natural extension that adapts k to the shape of each distribution—keeping more candidates when the model is uncertain, fewer when it is confident—and would be the next thing we would add.

5.7 Perplexity in Context

The final validation perplexity of **4.69** is worth interpreting against two reference points. A character-level bigram model on this corpus achieves roughly 3.4 perplexity, which is already surprisingly low because consecutive characters in English have very strong local correlations: given “th”, the next character is almost always “e”, “a”, or “o”. A fully-trained character-level LSTM on Shakespeare typically reaches ≈ 1.8 – 2.0 . Our model sits at 4.69, which is worse than the LSTM and—on a per-token basis—only modestly better than the bigram.

Why is the improvement over bigram so small relative to the model’s parameter count? Three factors explain this. First, the corpus is small: 1.1 M characters is roughly 10× smaller than

what LSTM baselines were tuned on. Second, the model is tiny: 0.82 M parameters against 5–10 M for a typical character LSTM. Third, training was brief: 5000 steps at batch size 64 is under 1.3 M training tokens seen—roughly 1.2 epochs over the corpus. The model has simply not been exposed to the data enough times to fit the higher-order structure that separates it from a bigram.

What this result *does* demonstrate, however, is that the architecture itself is correct. The model captures local structure (character statistics and short-range dependencies) exactly as expected, produces coherent short phrases, and generates text that is unambiguously Shakespeare-shaped rather than random. For a 4-minute training run on a laptop, that is the correct operating point. Closing the gap to the LSTM baseline would primarily require more training tokens and more epochs, not a change to the architecture.

5.8 Limitations

The main limitations of this work are:

1. **Corpus scale.** Tiny Shakespeare is small and low-entropy. Absolute perplexity values are therefore not directly comparable to those reported on datasets such as WikiText-103 or OpenWebText, where character-level models routinely reach perplexities in the 1.5–2.0 range.
2. **No KV-cache.** Generation currently recomputes the full context at every step, giving $O(T^2)$ per generated token instead of the $O(T)$ achievable with a KV-cache. For $T = 128$ this is imperceptible; for longer contexts it would dominate inference time.
3. **Hand-tuned hyperparameters.** The learning rate, batch size, dropout, and warmup length were chosen by hand based on common practice rather than through systematic search. A small grid or Bayesian sweep might find a better operating point, particularly for the interaction between peak learning rate and warmup length.
4. **Single seed.** All reported numbers come from a single training run. While the loss curve is smooth enough that seed-to-seed variance is unlikely to be large, a rigorous comparison against a baseline would average over multiple runs.

None of these limitations undermine the central claim of this report: that a decoder-only Transformer can be built from scratch, trained stably with modern optimisation techniques, and produce coherent autoregressive output—all within a few minutes on commodity hardware.

6 Conclusion

This report presented a complete, from-scratch implementation of a decoder-only Transformer (GPT-style) language model in PyTorch. Every core component—token embeddings, positional embeddings, causal multi-head self-attention, Pre-LN Transformer blocks, GELU feed-forward networks, and autoregressive generation—was built without relying on high-level wrappers such as `nn.TransformerDecoder`. This constraint turned out to be the most valuable part of the exercise: it forced us to reason about tensor shapes, masking semantics, and gradient flow at a level that high-level abstractions normally hide.

6.1 Summary of Findings

Training ran for 5000 steps and completed in **3.83 minutes** on a laptop GPU, converging to a final training loss of 1.613 and a validation loss of 1.545. In perplexity terms, the model achieves **4.69** on the held-out split—a meaningful improvement over the character-level bigram baseline of roughly 3.4 on this corpus, though still short of the ≈ 1.8 – 2.0 that a fully-trained character LSTM would reach. Given that the model has only 0.82 M parameters and sees roughly 1.2 epochs over the corpus, this operating point is exactly what the theory predicts: the model captures local character statistics and short-range syntax, but lacks the training exposure required to fit longer-range dependencies.

The three optimisation choices highlighted in the assignment all behaved as expected. Mixed precision (bf16 via `autocast`) preserved numerical stability throughout the run, with no NaNs, no loss spikes, and no evidence of gradient underflow. The linear-warmup-cosine-decay schedule produced the smoothest part of the training curve in its final phase, with validation loss descending monotonically from 1.559 at step 4250 to 1.545 at step 5000 as the learning rate decayed. Selective weight decay—applied to 18 tensors with $\dim(\theta) \geq 2$ but excluded from 34 one-dimensional tensors—matched the standard GPT recipe and required no further tuning.

The inference loop supports greedy decoding, temperature scaling, and top- k sampling. The comparison across five strategies from the same prompt was the most visually satisfying result of the project: greedy collapsed into a repetitive loop within fifteen tokens, temperature 0.8 with top- $k = 40$ produced fluent, varied Shakespeare-shaped text, and temperature 1.5 degenerated into gibberish. The coherence-diversity trade-off that the literature describes [2] showed up exactly as advertised, and being able to observe it directly—rather than take it on faith—was the point of building the model from scratch in the first place.

6.2 What This Project Taught Me

Beyond the mechanics, three things stand out.

First, the causal mask is not a minor implementation detail; it is the mechanism that makes

autoregressive training mathematically equivalent to parallel decoding. Until you have written the `masked_fill` call yourself and watched the lower-triangular attention heatmaps emerge, it is easy to treat the mask as boilerplate. It is not.

Second, mixed precision is not just a speed optimisation. Working through the interaction between `autocast`, `GradScaler`, gradient clipping, and unscaling is what clarified for me why `bf16` and `fp16` have different tooling requirements—and why getting the order of operations wrong silently destroys the clip threshold.

Third, and perhaps most importantly, the small implementation choices that textbooks often skim over—Pre-LN vs. Post-LN, whether to apply weight decay to LayerNorm gains, whether to tie the embedding and LM head weights—turn out to matter more than I expected for training stability and final model quality.

6.3 Future Work

Several extensions are natural next steps:

- **KV-cache.** The current generation loop recomputes the entire context at every step, giving $O(T^2)$ cost per generated token. A KV-cache would reduce this to $O(T)$ and make long-form generation practical.
- **Top- p (nucleus) sampling.** Adapting k to the shape of each token distribution would likely outperform fixed top- k , especially for prompts whose next-token entropy varies sharply.
- **Repetition penalties.** A small penalty on already-seen tokens would directly address the looping behaviour observed under greedy decoding.
- **Rotary Positional Embeddings (RoPE).** Replacing learned/sinusoidal embeddings with RoPE would scale better to longer contexts and is now standard in most modern open-weight LLMs [9].
- **Multi-GPU training.** Distributed Data Parallel or Fully Sharded Data Parallel would allow the model to grow past the single-GPU memory budget.
- **Parameter-efficient fine-tuning.** LoRA or a similar adapter-based method would make it possible to adapt the pretrained checkpoint to a different corpus without full retraining [3].

Looking back, the value of this assignment was not in producing state-of-the-art text—Tiny Shakespeare at 0.82 M parameters was never going to compete with a modern LLM. The value was in removing every layer of abstraction between the equations on the page and the tensors in memory. Having done that, I now read papers on Transformers differently: I can

see where the implementation complexity actually lives, and I understand which architectural choices are load-bearing versus which are cosmetic. For a 4-minute training run, that is a remarkably good return on investment.

References

- [1] T. Brown, B. Mann, N. Ryder, M. Subbiah, *et al.*, “Language models are few-shot learners,” *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [2] A. Holtzman, J. Buys, L. Du, M. Forbes, and Y. Choi, “The curious case of neural text degeneration,” *International Conference on Learning Representations (ICLR)*, 2020.
- [3] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, “Lora: Low-rank adaptation of large language models,” in *International Conference on Learning Representations (ICLR)*, 2022.
- [4] A. Karpathy, *Nanogpt: A minimal reproducible gpt implementation*, <https://github.com/karpathy/nanoGPT>, 2023.
- [5] I. Loshchilov and F. Hutter, “Sgdr: Stochastic gradient descent with warm restarts,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [6] I. Loshchilov and F. Hutter, “Decoupled weight decay regularization,” in *International Conference on Learning Representations (ICLR)*, 2019.
- [7] P. Micikevicius, S. Narang, J. Alben, G. Diamos, E. Elsen, D. Garcia, B. Ginsburg, M. Houston, O. Kuchaiev, G. Venkatesh, and H. Wu, “Mixed precision training,” in *International Conference on Learning Representations (ICLR)*, 2018.
- [8] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever, “Language models are unsupervised multitask learners,” OpenAI, Tech. Rep., 2019.
- [9] J. Su, Y. Lu, S. Pan, A. Murtadha, B. Wen, and Y. Liu, “Roformer: Enhanced transformer with rotary position embedding,” *arXiv preprint arXiv:2104.09864*, 2021.
- [10] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [11] R. Xiong, Y. Yang, D. He, K. Zheng, S. Zheng, C. Xing, H. Zhang, Y. Lan, L. Wang, and T.-Y. Liu, “On layer normalization in the transformer architecture,” in *International Conference on Machine Learning (ICML)*, 2020.